Nanostream Documentation

Release 0.1

Zachary Ernst

CONTENTS:

ONE

OVERVIEW

1.1 Why?

Tolstoy said that every happy family is the same, but every unhappy family is unhappy in its own way. ETL pipelines are unhappy families.

Why are they so unhappy? Every engineer who does more than one project involving ETL eventually goes through the same stages of ETL grief. First, they thing it's not so bad. Then they do another project and discover that they have to rewrite very similar code. Then they think, "Surely, I could have just written a few library functions and reused that code, saving lots of time." But when they try to do this, they discover that although their ETL projects are very similar, they are just different enough that their code isn't reusable. So they resign themselves to rewriting code over and over again. The code is unreliable, difficult to maintain, and usually poorly tested and documented because it's such a pain to write in the first place. The task of writing ETL pipelines is so lousy that engineering best practices tend to go out the window because the engineer has better things to do.

NanoStream is an ETL framework written in Python that tries to thread the ETL needle. It is opinionated about how ETL pipelines ought to be written; it provides a number of highly reusable and well-documented modules that are useful in ETL tasks; and it has just enough flexibility to allow an engineer to accommodate the strange idiosyncratic requirements of real-world ETL tasks.

The overall idea of NanoStream is simple. It is a streaming framework that resembles stream-processing patterns such as those found in Spark, Storm, and others. But unlike those tools, it requires no special infrastructure or server because it runs entirely within one process. It also differs from high-powered stream processing tools in that it is designed from the ground up with ETL in mind.

Despite its superficial similarity to Spark and Storm, it is not intended to compete with them in any way. NanoStream is not suitable for huge analytic pipelines that require massive amounts of computation. You won't want to analyze the Twitter firehose with NanoStream.

1.2 NanoStream pipelines

An ETL pipeline in NanoStream is a series of nodes connected by queues. Data is generated or processed in each node, and the output is placed on a queue to be picked up by downstream nodes.

For the sake of convenience, we distinguish between three types of nodes (although there's no real difference in their use or implementation):

1. Source nodes. These are nodes that generate data and send it to the rest of the pipeline. They might, for example, read data from an external data source such as an API endpoint or a database.

2. Worker nodes. The workers process data by picking up messages from their incoming queues. Their output is placed onto any number of outgoing queues to be further processed by downstream nodes.
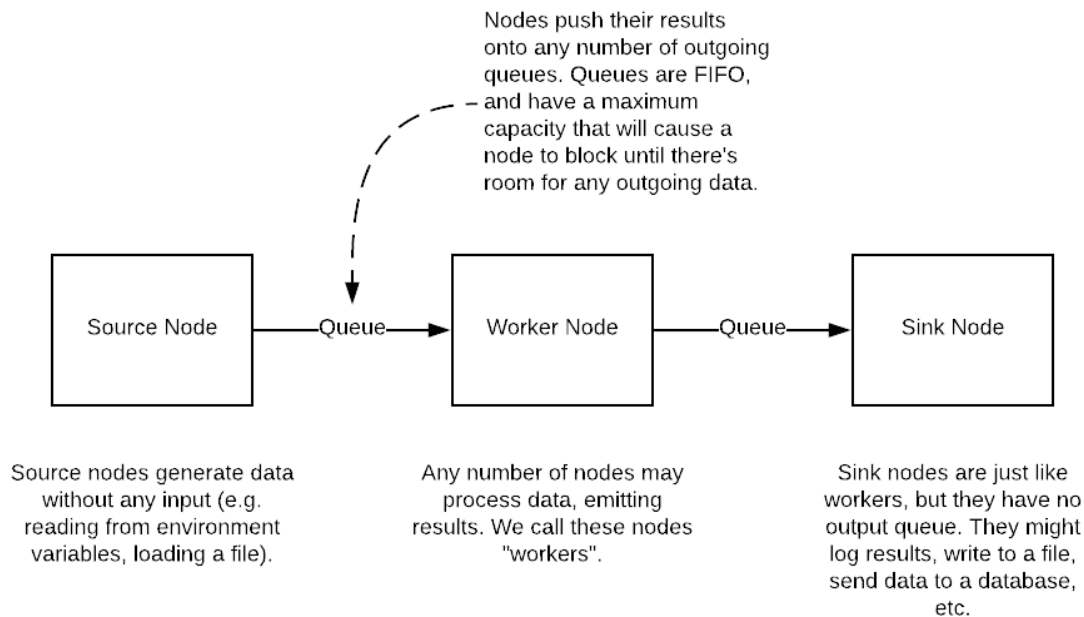
Fig. 1: Very high-level view of a NanoStream pipeline

3. Sink nodes. These are worker nodes with no outgoing queue. They will typically perform tasks such as inserting data into a database or generating statistics to be sent somewhere outside the pipeline.

All pipelines are implemented in pure Python (version >=3.5). Each node is instantiated from a class that inherits from the `NanoNode` class. Queues are never instantiated directly by the user; they are created automatically whenever two nodes are linked together.

There is a large (and growing) number of specialized `NanoNode` subclasses, each geared toward a specific task. Such tasks include:

1. Querying a table in a SQL database and sending the results downstream.

2. Making a request to a REST API, paging through the responses until there are no more results.

3. Ingesting individual messages from an upstream node and batching them together into a single message, or doing the reverse.

4. Reading environment variables.

5. Watching a directory for new files and sending the names of those files down the pipeline when they appear.

6. Filtering messages, letting them through the pipeline only if a particular test is passed.

All results and messages passed among the nodes must be dictionary-like objects. By default, messages hold the entire history of all the earlier messages that led to its being generated.

The goal is for NanoStream to be fully "batteries included", with built-in `NanoNode` subclasses for every necessary ETL task. But because this is actually impossible, we try to make it easy to roll your own node classes.
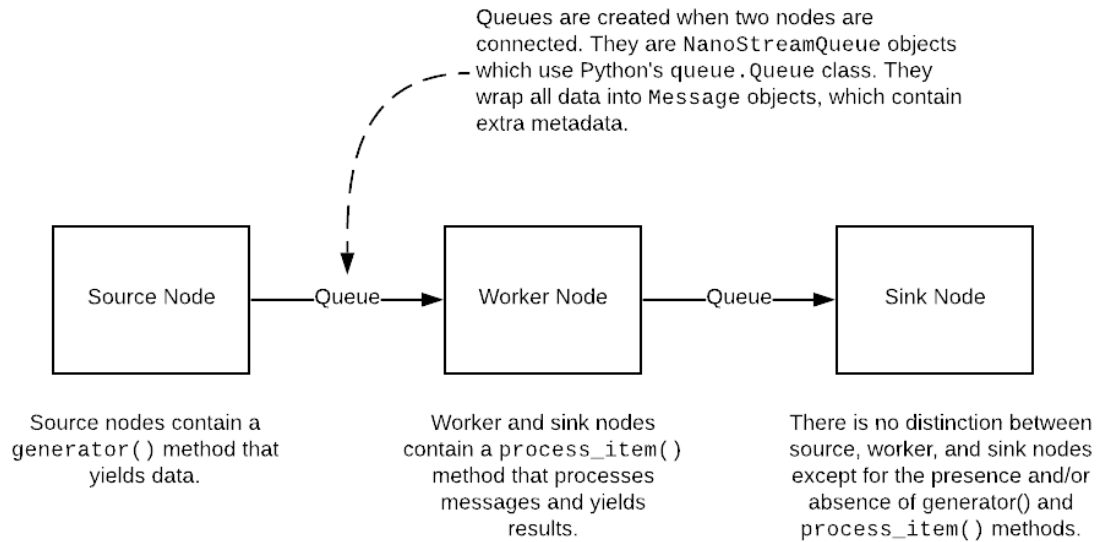
Queues are created when two nodes are
connected. They are `NanoStreamQueue` objects
which use Python's `queue.Queue` class. They
wrap all data into `Message` objects, which contain
extra metadata.

| Source Node | | Worker Node | | Sink Node |

Source nodes contain a
`generator()` method that
yields data.

Worker and sink nodes
contain a `process_item()`
method that processes
messages and yields
results.

There is no distinction between
source, worker, and sink nodes
except for the presence and/or
absence of generator() and
`process_item()` methods.

Fig. 2: Somewhat high-level view of a NanoStream pipeline

1.3 Using built-in `NanoNode` classes

The most straightforward use of NanoStream actually requires no coding, per se. You write a simple configuration file in YAML, and run the pipeline with the built-in command line tool, `nanostream_cli`. The tool reads the configuration file, instantiates the pipeline, and runs it.

The configuration file has two parts. The first specifies the nodes; the second specifies how they are linked together. Here is an example of a very simple configuration file:

```yaml
---
pipeline_name: Sample NanoStream configuration
pipeline_description: Reads some environment variables and prints them

nodes:
  get_environment_variables:
    class: GetEnvironmentVariables
    summary: Gets all the necessary environment variables
    options:
      environment_variables:
        - API_KEY
        - API_USER_ID

  print_variables:
    class: PrinterOfThings
    summary: Prints the environment variables to the terminal
    options:
      prepend: "Environment variables: "
```

```
paths:
  -

    - get_environment_variables
    - print_variables
```

Let's look at this config file one part at a time.

Other than the optional `pipeline_name` and `pipeline_description` fields, there are two top-level keys: `nodes` and `paths`. Each entry in the `node` section corresponds to a specific node in the pipeline. Their top-level key is whatever name you would like to use to refer to that node – it should be something short, descriptive, and Python-y. We recommend naming these as simple verb-noun phrases such as `get_environment_variable`, `print_variables`, `light_candle`, `curse_darkness`, etc.

There are three keys under each node configuration: `class`, `summary`, and `options`. The `class` key is the only one that's required. It is the name of the node's class. The `summary` key is just an optional arbitrary string.

All nodes share a certain number of common options that are important for any node. Depending on the node, there may also be one or more options specific to it. For example, a node that reads from a SQL database will likely have the table name as an option. In the above configuration, the `GetEnvironmentVariables` node has (reasonably enough) an option `environment_variables`, which contains a list of the environment variables that are to be retrieved. The `PrinterOfThings` node has an optional `prepend` value, which is a string that will be prepended to anything the node prints to the terminal.

The structure of the pipeline is given in the `paths` section, which contains a list of lists. Each list is a set of nodes that are to be linked together in order. In our example, the `paths` value says that `get_environment_variables` will send its output to `print_variables`. Paths can be arbitrarily long.

If you wanted to send the environment variables down two different execution paths, you add another list to the `paths`, like so:

```
paths:
  -

    - get_environment_variables
    - print_variables
  -

    - get_environment_variables
    - do_something_else
    - and_then_do_this
```

With this set of `paths`, the pipeline looks like a very simple tree, with `get_environment_variables` at the root, which branches to `print_variables` and `do_something_else`.

When you have written the configuration file, you're ready to use the NanoStream CLI. It accepts a command, followed by some options. As of now, the commands it accepts are `run`, which executes the pipeline, and `draw`, which generates a diagram of the pipeline. The relevant command(s) are:

```
python nanostream_cli.py [run | draw] --filename my_sample_config.yaml
```

It is also possible to skip using the configuration file and define your pipelines directly in code. In general, it's better to use the configuration file for a variety of reasons, but you always have the option of doing this in Python.

Nodes are defined in code by instantiating classes that inherit from `NanoNode`. Upon instantiation, the constructor takes the same set of keyword arguments as you see in the configuration. Nodes are linked together by the `>` operator, as in `node_1 > node_2`. After the pipeline has been built in this way, it is started by calling `node.global_start()` on any of the nodes in the pipeline.

The code corresponding to the configuration file above would look like this:

```
# Define the nodes using the various subclasses of NanoNode
get_environment_variables =
GetEnvironmentVariables(
    environment_variables=['API_KEY', 'API_USER_ID'])
print_variables = PrinterOfThings(prepend='Environment variables: ')

# The '>' operator can also be chained, as in:
# node_1 > node_2 > node_3 > ...
get_environment_variables > print_variables

# Run the pipeline. This command will not block.
get_environment_variables.global_start()
```

1.4 Rolling your own `NanoNode` class

If there are no built-in `NanoNode` classes suitable for your ETL pipeline, it is easy to write your own.

For example, suppose you want to create a source node for your pipeline that simply emits a user-defined string every few seconds forever. The user would be able to specify the string and the number of seconds to pause after each message has been sent. The class could be defined like so:

```
class FooEmitter(NanoNode):  # inherit from NanoNode
    '''
    Sends ``self.output_string`` every ``self.interval`` seconds.
    '''
    def __init__(self, output_string='', interval=1, **kwargs):
        self.output_string = output_string
        self.interval = interval
        super(FooEmitter, self).__init__()  # Must call the `NanoNode` __init__

    def generator(self):
        while True:
            time.sleep(self.interval)
            yield self.output_string  # Output must be yielded, not returned
```

Let's look at each part of this class.

The first thing to note is that the class inherits from `NanoNode` – this is the mix-in class that gives the node all of its functionality within the NanoStream framework.

The __init__ method should take only keyword arguments, not positional arguments. This restriction is to guarantee that the configuration files have names for any options that are specified in the pipeline. In the __init__ function, you should also be sure to accept `**kwargs`, because options that are common to all `NanoNode` objects are expected to be there.

After any attributes have been defined, the __init__ method **must** invoke the parent class's constructor through the use of the `super` function. Be sure to pass the `**kwargs` argument into the function as shown in the example.

If the node class is intended to be used as a source node, then you need to define a `generator` method. This method can be virtually anything, so long as it sends its output via a `yield` statement.

If you need to define a worker node (that is, a node that accepts input from a queue), you will provide a `process_item` method instead of a `generator`. But the structure of that method is the same, with the single exception that you will have access to a __message__ attribute which contains the incoming message data. The structure of a typical `process_item` method is shown in the figure.
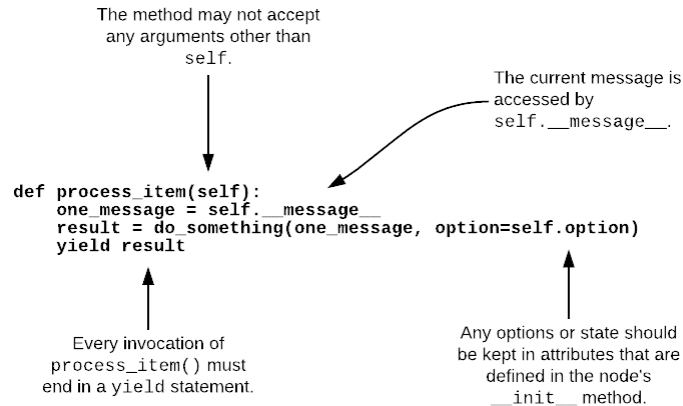
Fig. 3: A typical `process_item` method for `NanoNode` objects

For example, let's suppose you want to create a node that is passed a string as a message, and returns `True` if the message has an even number of characters, `False` otherwise. The class definition would look like this:

```python
class MessageLengthTester(NanoNode):
    def __init__(self):
        # No particular initialization required in this example
        super(MessageLengthTester, self).__init__()

    def process_item(self):
        if len(self.__message__) % 2 == 0:
            yield True
        else:
            yield False
```

1.5 Composing and configuring `NanoNode` objects

> **Warning:** The code described in this section is experimental and very unstable. It would be bad to use it for anything important.

Let's suppose you've worked very hard to create the pipeline from the last example. Now, your boss says that another engineering team wants to use it, but they want to rename parameters and "freeze" the values of certain other parameters to specific values. Once that's done, they want to use it as just one part of a more complicated `NanoStream` pipeline.

This can be accomplished using a configuration file. When `NanoStream` parses the configuration file, it will dynamically create the desired class, which can be instantiated and used as if it were a single node in another pipeline.

The configuration file is written in YAML, and it would look like this:

```yaml
name: FooMessageTester

nodes:
```

```
- name: foo_generator
  class FooEmitter
  frozen_arguments:
    message: foobar
  arg_mapping:
    interval: foo_interval
- name: length_tester
  class: MessageLengthTester
  arg_mapping: null
```

With this file saved as (e.g.) `foo_message.yaml`, the following code will create a `FooMessageTester` class and instantiate it:

```
foo_message_config = yaml.load(open('./foo_message.yaml', 'r').read())
class_factory(foo_message_config)
# At this point, there is now a `FooMessageTester` class
foo_node = FooMessageTester(foo_interval=1)
```

You can now use `foo_node` just as you would any other node. So in order to run it, you just do:

```
foo_node.global_start()
```

Because `foo_node` is just another node, you can insert it into a larger pipeline and reuse it. For example, suppose that other engineering team wants to add a `PrinterOfThings` to the end of the pipeline. They'd do this:

```
printer = PrinterOfThings()
foo_node > printer
```

TWO

THE DATA JOURNEY

2.1 Overview

NanoStream pipelines create dictionary-like objects as messages, and those messages move through the various nodes until they reach a sink. As they move through the nodes, they are modified in one or more of the following ways:

1. New keys and values are added to the dictionary.

2. Keys and values are removed from the dictionary.

3. Values are modified through in-place operations.

4. The structure of the dictionary is changed (e.g. keys are renamed, the dictionary is flattened, and so on).

Which of these operations is used depends on the particular type of node that processes the message, and how the various options are specified for that node.

By default, each time a message is processed, all of its existing keys and values are retained in the message as it's passed to the next node. This behavior is to enable the message to accumulate information over several steps in the pipeline because some nodes require data that is generated by various operations.

Nodes may have access to the entire message, or it is possible to specify which key-value pair is passed to it. This is done by using the `key` option in the node definition. If the node will be generating results to be passed downstream, then we need to either (1) specify the for those results; or (2) make sure that the node is generating a dictionary. If (2), then by default the dictionary will be merged into the incoming message, and the combined dictionary will be placed on the node's outgoing queue. If a specific key is specified for the generated data, then we use the `output_key` option in the node definition.

2.2 Example: Making a GET request

Let's consider a very common ETL task. We want to make a GET request to an API endpoint and return the result. The GET request will take a couple of parameters, such as an endpoint name, date and username. For our example, the URL will just be:

```
http://example.api.com/ENDPOINT?date=DATE&username=USERNAME
```

As you would expect, the username and endpoint can be specified in advance; but the date will change each time the pipeline is run. So the date has to be generated at runtime and passed to the node. In our example, the date will be passed to the pipeline as an environment variables `DATE` when the pipeline is executed.

NanoNode contains node classes for reading environment variables and for making GET requests. They are called `GetEnvironmentVariables` and `HttpGetRequest`. They take the following options:

1. `GetEnvironmentVariables`

- • environment_variables: A list of the names of the environment variables to be fetched. The results will be put in keys named after those environment variables.

2. HttpGetRequest

  - • url (required): The URL for the GET request. Any parameters that will be filled-in at runtime should be put into curly braces. See the example configuration file below.

  - • json (optional: default True) Whether the response should be parsed as JSON.

  - • endpoint_dict (optional) Keys and values to be substituted into the url.

THREE

REQUIREMENTS

NanoStream is written in Python 3.5. All requirements are pip-installable and are listed n the `requirements.txt` file:

` pip install -r requirements.txt `

The documentation is written using Sphinx, so if you want to rebuild the docs, you'll do:

` make [html | latexpdf | whatever] `

That ought to be everything.

IMPLEMENTATION

This section describes what's happening under the hood in a `NanoStream` data pipeline. Most people won't need to read this section.

4.1 The data journey

NanoStream pipelines are sets of `NanoNode` objects connected by `NanoStreamQueue` objects. Think of each `NanoNode` as a vertex in a directed graph, and each `NanoStreamQueue` as a directed edge.

There are two types of `NanoNode` objects. A "source" is a `NanoNode` that does not accept incoming data from another `NanoNode`. A "processor" is any `NanoNode` that is not a "source". Note that there is nothing in the class definition or object that distinguishes between these two – the only difference is that processors have a `process_item` method, and sources have a `generator` method. Other than that, they are identical.

The data journey begins with one or more source nodes. When a source node is started (by calling its `start` method), a new thread is created and the node's `generator` method is executed inside the thread. As results from the `generator` method are yielded, they are placed on each outgoing `NanoStreamQueue` to be picked up by one or more processors downstream.

The data from the source's `generator` is handled by the `NanoStreamQueue` object. At its heart, the `NanoStreamQueue` is simply a class which has a Python `Queue.queue` object as an attribute. The reason we don't simply use Python `Queue` objects is because the `NanoStreamQueue` contains some logic that's useful. In particular:

1. It wraps the data into a `NanoStreamMessage` object, which also holds useful metadata including a UUID, the ID of the node that generated the data, and a timestamp.

2. If the `NanoStreamQueue` receives data that is simply a `None` object, then it is skipped.

API DOCUMENTATION

5.1 Node module

The `node` module contains the `NanoNode` class, which is the foundation for NanoStream.

class nanostream.node.**AggregateValues**(*values=False*, *tail_path=None*, ***kwargs*)
>   Bases: *nanostream.node.NanoNode*

>   Does that.

>   **process_item**()
>>   Default no-op for nodes.

class nanostream.node.**BatchMessages**(*batch_size=None*, *batch_list=None*, *counter=0*, *time-out=5*, ***kwargs*)
>   Bases: *nanostream.node.NanoNode*

>   **cleanup**()

>   **process_item**()
>>   Default no-op for nodes.

class nanostream.node.**CSVReader**(**args*, ***kwargs*)
>   Bases: *nanostream.node.NanoNode*

>   **process_item**()
>>   Default no-op for nodes.

class nanostream.node.**CSVToDictionaryList**(***kwargs*)
>   Bases: *nanostream.node.NanoNode*

>   **process_item**()
>>   Default no-op for nodes.

class nanostream.node.**ConstantEmitter**(*thing=None*, *delay=2*, ***kwargs*)
>   Bases: *nanostream.node.NanoNode*

>   Send a thing every n seconds

>   **generator**()

class nanostream.node.**CounterOfThings**(**args*, *batch=False*, *get_runtime_attrs=<function no_op>*, *get_runtime_attrs_args=None*, *get_runtime_attrs_kwargs=None*, *runtime_attrs_destinations=None*, *input_mapping=None*, *retain_input=True*, *throttle=0*, *keep_alive=True*, *max_errors=0*, *name=None*, *input_message_keypath=None*, *key=None*, *messages_received_counter=0*, *messages_sent_counter=0*, *post_process_function=None*, *post_process_keypath=None*, *summary=''*, *post_process_function_kwargs=None*, *output_key=None*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> **foo__init__**(**args*, *start=0*, *end=None*, ***kwargs*)

> **generator**()
> > Just start counting integers

class nanostream.node.**DynamicClassMediator**(**args*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> **get_sink**()

> **get_source**()

> **hi**()

> **sink_list**()

> **source_list**()

class nanostream.node.**Filter**(*test=None*, *key=None*, *value=True*, **args*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> Applies tests to each message and filters out messages that don't pass

> **Built-in tests:** key_exists value_is_true value_is_not_none

> **Example:**

> > {'test': 'key_exists', 'key': mykey}

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**GetEnvironmentVariables**(*mappings=None*, *environment_variables=None*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> **generator**()

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**InsertData**(*overwrite=True*, *overwrite_if_null=True*, *value_dict=None*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**LocalDirectoryWatchdog**(*directory='.'*, *check_interval=3*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> **generator**()

class nanostream.node.**LocalFileReader**(**args*, ***kwargs*)
> Bases: *nanostream.node.NanoNode*
>
> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**NanoNode**(**args*, *batch=False*, *get_runtime_attrs=<function no_op>*, *get_runtime_attrs_args=None*, *get_runtime_attrs_kwargs=None*, *runtime_attrs_destinations=None*, *input_mapping=None*, *retain_input=True*, *throttle=0*, *keep_alive=True*, *max_errors=0*, *name=None*, *input_message_keypath=None*, *key=None*, *messages_received_counter=0*, *messages_sent_counter=0*, *post_process_function=None*, *post_process_keypath=None*, *summary=''*, *post_process_function_kwargs=None*, *output_key=None*, ***kwargs*)
> Bases: object
>
> The foundational class of *NanoStream*. This class is inherited by all nodes in a computation graph.
>
> Order of operations: 1. Child class __init__ function 2. NanoNode __init__ function 3. preflight_function (Specified in initialization params) 4. setup 5. start
>
> These methods have the following intended uses:
>
> 1. __init__ Sets attribute values and calls the NanoNode __init__ method.
>
> 2. get_runtime_attrs Sets any attribute values that are to be determined at runtime, e.g. by checking environment variables or reading values from a database. The get_runtime_attrs should return a dictionary of attributes -> values, or else None.
>
> 3. setup Sets the state of the NanoNode and/or creates any attributes that require information available only at runtime.
>
> > **Variables**
> >
> > - **send_batch_markers** – If True, then a BatchStart marker will be sent when a new input is received, and a BatchEnd will be sent after the input has been processed. The intention is that a number of items will be emitted for each input received. For example, we might emit a table row-by-row for each input.
> >
> > - **get_runtime_attrs** – A function that returns a dictionary-like object. The keys and values will be saved to this NanoNode object's attributes. The function is executed one time, upon starting the node.
> >
> > - **get_runtime_attrs_args** – A tuple of arguments to be passed to the get_runtime_attrs function upon starting the node.
> >
> > - **get_runtime_attrs_kwargs** – A dictionary of kwargs passed to the get_runtime_attrs function.
> >
> > - **runtime_attrs_destinations** – If set, this is a dictionary mapping the keys returned from the get_runtime_attrs function to the names of the attributes to which the values will be saved.
> >
> > - **throttle** – For each input received, a delay of throttle seconds will be added.
> >
> > - **keep_alive** – If True, keep the node's thread alive after everything has been processed.
> >
> > - **name** – The name of the node. Defaults to a randomly generated hash. Note that this hash is not consistent from one run to the next.

- **input_mapping** – When the node receives a dictionary-like object, this dictionary will cause the keys of the dictionary to be remapped to new keys.

- **retain_input** – If True, then combine the dictionary-like input with the output. If keys clash, the output value will be kept.

- **input_message_keypath** – Read the value in this keypath as the content of the incoming message.

add_edge(*target*, ***kwargs*)
Create an edge connecting *self* to *target*. The edge is really just a queue

all_connected(*seen=None*)
Returns all the nodes connected (directly or indirectly) to self.

Args:

seen (set): A set of all the nodes that have been identified as connected to self.

Returns:

(set of **NanoNode**): All the nodes connected to **self**. This includes self.

broadcast(*broadcast_message*)
Puts the message into all the input queues for all connected nodes.

cleanup()

draw_pipeline()
Draw the pipeline structure using graphviz.

global_start(*datadog=False*, *pipeline_name=None*)
Starts every node connected to self.

input_queue_size
Return the total number of items in all of the queues that are inputs to this node.

is_sink
Tests whether the node is a sink or not, i.e. whether there are no outputs from the node.

is_source
Tests whether the node is a source or not, i.e. whether there are no inputs to the node.

kill_pipeline()

log_info(*message=''*)

process_item(**args*, ***kwargs*)
Default no-op for nodes.

processor()
This calls the user's process_item with just the message content, and then returns the full message.

setup()
To be overridden by child classes when we need to do something after setting attributes and the pre-flight function.

start()
Starts the node. This is called by NanoNode.global_start().

stream()
Called in each NanoNode thread.

terminate_pipeline()
This method can be called on any node in a pipeline, and it will cause all of the nodes to terminate.

thread_monitor()
> This function loops over all of the threads in the pipeline, checking that they are either `finished` or `running`. If any have had an abnormal exit, terminate the entire pipeline.

time_running
> Return the number of wall-clock seconds elapsed since the node was started.

class nanostream.node.**NothingToSeeHere**
> Bases: `object`

> Vacuous class used as a no-op message type.

class nanostream.node.**Parameters**(***kwargs*)
> Bases: `object`

class nanostream.node.**PrinterOfThings**(**args*, ***kwargs*)
> Bases: *[nanostream.node.NanoNode](#)*

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**RandomSample**(*sample=0.1*)
> Bases: *[nanostream.node.NanoNode](#)*

> Lets through only a random sample of incoming messages. Might be useful for testing, or when only approximate results are necessary.

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**Remapper**(*mapping=None*, ***kwargs*)
> Bases: *[nanostream.node.NanoNode](#)*

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**SequenceEmitter**(*sequence*, **args*, *max_sequences=1*, ***kwargs*)
> Bases: *[nanostream.node.NanoNode](#)*

> Emits `sequence` max_sequences times, or forever if `max_sequences` is `None`.

> **generator**()
> > Emit the sequence max_sequences times.

> **process_item**()
> > Emit the sequence max_sequences times.

class nanostream.node.**Serializer**(*values=False*, **args*, ***kwargs*)
> Bases: *[nanostream.node.NanoNode](#)*

> Takes an iterable thing as input, and successively yields its items.

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**SimpleTransforms**(*missing_keypath_action='ignore'*, *starting_path=None*, *transform_mapping=None*, *target_value=None*, *keypath=None*, ***kwargs*)
> Bases: *[nanostream.node.NanoNode](#)*

> **process_item**()
> > Default no-op for nodes.

class nanostream.node.**StreamMySQLTable**(**args*, *host='localhost'*, *user=None*, *table=None*, *password=None*, *database=None*, *port=3306*, *to_row_obj=False*, *send_batch_markers=True*, ***kwargs*)

Bases: *nanostream.node.NanoNode*

generator()

get_schema()

setup()
    To be overridden by child classes when we need to do something after setting attributes and the pre-flight function.

class nanostream.node.**StreamingJoin**(*window=30*, *streams=None*, **args*, ***kwargs*)
    Bases: *nanostream.node.NanoNode*

Joins two streams on a key, using exact match only. MVP.

process_item()

class nanostream.node.**SubstituteRegex**(*match_regex=None*, *substitute_string=None*, **args*, ***kwargs*)

Bases: *nanostream.node.NanoNode*

process_item()
    Default no-op for nodes.

class nanostream.node.**TimeWindowAccumulator**(**args*, ***kwargs*)
    Bases: *nanostream.node.NanoNode*

Every N seconds, put the latest M seconds data on the queue.

class nanostream.node.**bcolors**
    Bases: object

This class holds the values for the various colors that are used in the tables that monitor the status of the nodes.

BOLD = '\x1b[1m'

ENDC = '\x1b[0m'

FAIL = '\x1b[91m'

HEADER = '\x1b[95m'

OKBLUE = '\x1b[94m'

OKGREEN = '\x1b[92m'

UNDERLINE = '\x1b[4m'

WARNING = '\x1b[93m'

nanostream.node.**class_factory**(*raw_config*)

nanostream.node.**get_environment_variables**(**args*)
    Retrieves the environment variables listed in `*args`.

    Args: args (list of str): List of environment variables.

    Returns:

        dict: Dictionary of environment variables to values. If the environment variable is not defined, the value is `None`.

nanostream.node.**get_node_dict**(*node_config*)

---

nanostream.node.**kwarg_remapper**(*f*, ***kwarg_mapping*)

nanostream.node.**no_op**(**args*, ***kwargs*)
    No-op function to serve as default get_runtime_attrs.

nanostream.node.**template_class**(*class_name*,   *parent_class*,   *kwargs_remapping*,
                            frozen_arguments_mapping)

5.2 Data structures module

Data types (e.g. Rows, Records) for ETL.

class nanostream.utils.data_structures.**BOOL**(*value*, *original_type=None*, *name=None*)
    Bases: *nanostream.utils.data_structures.DataType*,   *nanostream.utils.*
    data_structures.IntermediateTypeSystem

    python_cast_function
        alias of builtins.bool

class nanostream.utils.data_structures.**DATETIME**(*value*,   *original_type=None*,
                            name=None)
    Bases: *nanostream.utils.data_structures.DataType*,   *nanostream.utils.*
    data_structures.IntermediateTypeSystem

    python_cast_function()

class nanostream.utils.data_structures.**DataSourceTypeSystem**
    Bases: object

    Information about mapping one type system onto another contained in the children of this class.

    static convert(*obj*)
        Override this method if something more complicated is necessary.

    static type_mapping(**args*, ***kwargs*)

class nanostream.utils.data_structures.**DataType**(*value*,   *original_type=None*,
                            name=None)
    Bases: object

    Each DataType gets a python_cast_function, which is a function.

    intermediate_type = None

    python_cast_function = None

    to_intermediate_type()
        Convert the DataType to an IntermediateDataType using its class's intermediate_type
        attribute.

    to_python()

    type_system
        Just for convenience to make the type system an attribute.

class nanostream.utils.data_structures.**FLOAT**(*value*, *original_type=None*, *name=None*)
    Bases: *nanostream.utils.data_structures.DataType*,   *nanostream.utils.*
    data_structures.IntermediateTypeSystem

    python_cast_function
        alias of builtins.float

class nanostream.utils.data_structures.**INTEGER**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.DataType](#)*, *[nanostream.utils.data_structures.IntermediateTypeSystem](#)*

python_cast_function
    alias of builtins.int

exception nanostream.utils.data_structures.**IncompatibleTypesException**

Bases: Exception

class nanostream.utils.data_structures.**IntermediateTypeSystem**

Bases: *[nanostream.utils.data_structures.DataSourceTypeSystem](#)*

Never instantiate this by hand.

class nanostream.utils.data_structures.**MYSQL_BOOL**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.DataType](#)*, *[nanostream.utils.data_structures.MySQLTypeSystem](#)*

intermediate_type
    alias of *[BOOL](#)*

python_cast_function
    alias of builtins.bool

class nanostream.utils.data_structures.**MYSQL_DATE**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.DataType](#)*, *[nanostream.utils.data_structures.MySQLTypeSystem](#)*

intermediate_type
    alias of *[DATETIME](#)*

python_cast_function()

class nanostream.utils.data_structures.**MYSQL_ENUM**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.DataType](#)*, *[nanostream.utils.data_structures.MySQLTypeSystem](#)*

intermediate_type
    alias of *[STRING](#)*

python_cast_function
    alias of builtins.str

class nanostream.utils.data_structures.**MYSQL_INTEGER**

Bases: type

class nanostream.utils.data_structures.**MYSQL_INTEGER0**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

max_length = 0

class nanostream.utils.data_structures.**MYSQL_INTEGER1**(*value*, *original_type=None*, *name=None*)

Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

max_length = 1

class nanostream.utils.data_structures.**MYSQL_INTEGER10**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 10

class nanostream.utils.data_structures.**MYSQL_INTEGER1024**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 1024

class nanostream.utils.data_structures.**MYSQL_INTEGER11**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 11

class nanostream.utils.data_structures.**MYSQL_INTEGER12**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 12

class nanostream.utils.data_structures.**MYSQL_INTEGER128**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 128

class nanostream.utils.data_structures.**MYSQL_INTEGER13**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 13

class nanostream.utils.data_structures.**MYSQL_INTEGER14**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 14

class nanostream.utils.data_structures.**MYSQL_INTEGER15**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 15

class nanostream.utils.data_structures.**MYSQL_INTEGER16**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 16

class nanostream.utils.data_structures.**MYSQL_INTEGER16384**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

    max_length = 16384

class nanostream.utils.data_structures.**MYSQL_INTEGER17**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

> **max_length = 17**

class nanostream.utils.data_structures.**MYSQL_INTEGER18**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 18**

class nanostream.utils.data_structures.**MYSQL_INTEGER19**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 19**

class nanostream.utils.data_structures.**MYSQL_INTEGER2**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 2**

class nanostream.utils.data_structures.**MYSQL_INTEGER20**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 20**

class nanostream.utils.data_structures.**MYSQL_INTEGER2048**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 2048**

class nanostream.utils.data_structures.**MYSQL_INTEGER21**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 21**

class nanostream.utils.data_structures.**MYSQL_INTEGER22**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 22**

class nanostream.utils.data_structures.**MYSQL_INTEGER23**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 23**

class nanostream.utils.data_structures.**MYSQL_INTEGER24**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 24**

class nanostream.utils.data_structures.**MYSQL_INTEGER25**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

> **max_length = 25**

class nanostream.utils.data_structures.**MYSQL_INTEGER256**(*value*, *original_type=None*,
  name=None)
  Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

max_length = 256

class nanostream.utils.data_structures.**MYSQL_INTEGER26**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 26

class nanostream.utils.data_structures.**MYSQL_INTEGER27**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 27

class nanostream.utils.data_structures.**MYSQL_INTEGER28**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 28

class nanostream.utils.data_structures.**MYSQL_INTEGER29**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 29

class nanostream.utils.data_structures.**MYSQL_INTEGER3**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 3

class nanostream.utils.data_structures.**MYSQL_INTEGER30**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 30

class nanostream.utils.data_structures.**MYSQL_INTEGER31**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 31

class nanostream.utils.data_structures.**MYSQL_INTEGER32**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 32

class nanostream.utils.data_structures.**MYSQL_INTEGER32768**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 32768

class nanostream.utils.data_structures.**MYSQL_INTEGER4**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

    max_length = 4

class nanostream.utils.data_structures.**MYSQL_INTEGER4096**(*value*, *original_type=None*,
name=None)

    Bases: *nanostream.utils.data_structures.MYSQL_INTEGER_BASE*

```
max_length = 4096
```

class nanostream.utils.data_structures.**MYSQL_INTEGER5**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 5
```

class nanostream.utils.data_structures.**MYSQL_INTEGER512**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 512
```

class nanostream.utils.data_structures.**MYSQL_INTEGER6**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 6
```

class nanostream.utils.data_structures.**MYSQL_INTEGER64**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 64
```

class nanostream.utils.data_structures.**MYSQL_INTEGER7**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 7
```

class nanostream.utils.data_structures.**MYSQL_INTEGER8**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 8
```

class nanostream.utils.data_structures.**MYSQL_INTEGER8192**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 8192
```

class nanostream.utils.data_structures.**MYSQL_INTEGER9**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_INTEGER_BASE](#)*

```
max_length = 9
```

class nanostream.utils.data_structures.**MYSQL_INTEGER_BASE**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.DataType](#)*, *[nanostream.utils.data_structures.MySQLTypeSystem](#)*

    intermediate_type
        alias of *[INTEGER](#)*

    python_cast_function
        alias of `builtins.int`

class nanostream.utils.data_structures.**MYSQL_VARCHAR**

    Bases: `type`

class nanostream.utils.data_structures.**MYSQL_VARCHAR0**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 0

class nanostream.utils.data_structures.**MYSQL_VARCHAR1**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 1

class nanostream.utils.data_structures.**MYSQL_VARCHAR10**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 10

class nanostream.utils.data_structures.**MYSQL_VARCHAR1024**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 1024

class nanostream.utils.data_structures.**MYSQL_VARCHAR11**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 11

class nanostream.utils.data_structures.**MYSQL_VARCHAR12**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 12

class nanostream.utils.data_structures.**MYSQL_VARCHAR128**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 128

class nanostream.utils.data_structures.**MYSQL_VARCHAR13**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 13

class nanostream.utils.data_structures.**MYSQL_VARCHAR14**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 14

class nanostream.utils.data_structures.**MYSQL_VARCHAR15**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 15

class nanostream.utils.data_structures.**MYSQL_VARCHAR16**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 16

class nanostream.utils.data_structures.**MYSQL_VARCHAR16384**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 16384

class nanostream.utils.data_structures.**MYSQL_VARCHAR17**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 17

class nanostream.utils.data_structures.**MYSQL_VARCHAR18**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 18

class nanostream.utils.data_structures.**MYSQL_VARCHAR19**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 19

class nanostream.utils.data_structures.**MYSQL_VARCHAR2**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 2

class nanostream.utils.data_structures.**MYSQL_VARCHAR20**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 20

class nanostream.utils.data_structures.**MYSQL_VARCHAR2048**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 2048

class nanostream.utils.data_structures.**MYSQL_VARCHAR21**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 21

class nanostream.utils.data_structures.**MYSQL_VARCHAR22**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 22

class nanostream.utils.data_structures.**MYSQL_VARCHAR23**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 23

class nanostream.utils.data_structures.**MYSQL_VARCHAR24**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 24

class nanostream.utils.data_structures.**MYSQL_VARCHAR25**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 25

class nanostream.utils.data_structures.**MYSQL_VARCHAR256**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 256

class nanostream.utils.data_structures.**MYSQL_VARCHAR26**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 26

class nanostream.utils.data_structures.**MYSQL_VARCHAR27**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 27

class nanostream.utils.data_structures.**MYSQL_VARCHAR28**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 28

class nanostream.utils.data_structures.**MYSQL_VARCHAR29**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 29

class nanostream.utils.data_structures.**MYSQL_VARCHAR3**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 3

class nanostream.utils.data_structures.**MYSQL_VARCHAR30**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 30

class nanostream.utils.data_structures.**MYSQL_VARCHAR31**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 31

class nanostream.utils.data_structures.**MYSQL_VARCHAR32**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 32

class nanostream.utils.data_structures.**MYSQL_VARCHAR32768**(*value*, *original_type=None*, *name=None*)

    Bases: *[nanostream.utils.data_structures.MYSQL_VARCHAR_BASE](#)*

    max_length = 32768

class nanostream.utils.data_structures.**MYSQL_VARCHAR4**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 4

class nanostream.utils.data_structures.**MYSQL_VARCHAR4096**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 4096

class nanostream.utils.data_structures.**MYSQL_VARCHAR5**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 5

class nanostream.utils.data_structures.**MYSQL_VARCHAR512**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 512

class nanostream.utils.data_structures.**MYSQL_VARCHAR6**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 6

class nanostream.utils.data_structures.**MYSQL_VARCHAR64**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 64

class nanostream.utils.data_structures.**MYSQL_VARCHAR7**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 7

class nanostream.utils.data_structures.**MYSQL_VARCHAR8**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 8

class nanostream.utils.data_structures.**MYSQL_VARCHAR8192**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 8192

class nanostream.utils.data_structures.**MYSQL_VARCHAR9**(*value*, *original_type=None*, *name=None*)

    Bases: *nanostream.utils.data_structures.MYSQL_VARCHAR_BASE*

    max_length = 9

class nanostream.utils.data_structures.**MYSQL_VARCHAR_BASE**(*value*, *original_type=None*, *name=None*)

Bases: [*nanostream.utils.data_structures.DataType*](), [*nanostream.utils.data_structures.MySQLTypeSystem*]()

intermediate_type
: alias of [*STRING*]()

python_cast_function
: alias of `builtins.str`

class nanostream.utils.data_structures.**MySQLTypeSystem**
: Bases: [*nanostream.utils.data_structures.DataSourceTypeSystem*]()

Each `TypeSystem` gets a `type_mapping` static method that takes a string and returns the class in the type system named by that string. For example, `int(8)` in a MySQL schema should return the `MYSQL_INTEGER8` class.

static type_mapping(*string*)
: Parses the schema strings from MySQL and returns the appropriate class.

class nanostream.utils.data_structures.**PrimitiveTypeSystem**
: Bases: [*nanostream.utils.data_structures.DataSourceTypeSystem*]()

class nanostream.utils.data_structures.**PythonTypeSystem**
: Bases: [*nanostream.utils.data_structures.DataSourceTypeSystem*]()

class nanostream.utils.data_structures.**Row**(**records*, *type_system=None*)
: Bases: `object`

A collection of `DataType` objects (typed values). They are dictionaries mapping the names of the values to the `DataType` objects.

concat(*other*, *fail_on_duplicate=True*)

static from_dict(*row_dictionary*, ***kwargs*)
: Creates a `Row` object form a dictionary mapping names to values.

is_empty()

keys()
: For implementing the mapping protocol.

class nanostream.utils.data_structures.**STRING**(*value*, *original_type=None*, *name=None*)
: Bases: [*nanostream.utils.data_structures.DataType*](), [*nanostream.utils.data_structures.IntermediateTypeSystem*]()

python_cast_function
: alias of `builtins.str`

nanostream.utils.data_structures.**all_bases**(*obj*)
: Return all the class to which `obj` belongs.

nanostream.utils.data_structures.**convert_to_type_system**(*obj*, *cls*)

nanostream.utils.data_structures.**get_type_system**(*obj*)

nanostream.utils.data_structures.**make_types**()

nanostream.utils.data_structures.**mysql_type**(*string*)
: Parses the schema strings from MySQL and returns the appropriate class.

nanostream.utils.data_structures.**primitive_to_intermediate_type**(*thing*, *name=None*)

5.3 Network nodes module

Classes that deal with sending and receiving data across the interwebs.

class nanostream.node_classes.network_nodes.**HttpGetRequest**(*url=None*, *endpoint_dict=None*, *json=True*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> Node class for making simple GET requests.

> **process_item**()
>> The input to this function will be a dictionary-like object with parameters to be substituted into the endpoint string and a dictionary with keys and values to be passed in the GET request.

>> Three use-cases: 1. Endpoint and parameters set initially and never changed. 2. Endpoint and parameters set once at runtime 3. Endpoint and parameters set by upstream messages

class nanostream.node_classes.network_nodes.**HttpGetRequestPaginator**(*endpoint_dict=None*, *json=True*, *pagination_get_request_key=None*, *endpoint_template=None*, *additional_data_key=None*, *pagination_key=None*, *pagination_template_key=None*, *default_offset_value=''*, ***kwargs*)

> Bases: *nanostream.node.NanoNode*

> Node class for HTTP API requests that require paging through sets of results.

> **process_item**()
>> Default no-op for nodes.

class nanostream.node_classes.network_nodes.**PaginatedHttpGetRequest**(*endpoint_template=None*, *additional_data_key=None*, *pagination_key=None*, *pagination_get_request_key=None*, *default_offset_value=''*, *additional_data_test=<class 'bool'>*)

> Bases: object

> For handling requests in a semi-general way that require paging through lists of results and repeatedly making GET requests.

responses()
>> Generator. Yields each response until empty.

5.4 NanoStreamMessage module

The `NanoStreamMesaage` encapsulates the content of each piece of data, along with some useful metadata.

class `nanostream.message.message.`**NanoStreamMessage**(*message_content*)
> Bases: `object`

> A class that contains the message payloads that are queued for each `NanoStreamProcessor`. It holds the messages and lots of metadata used for logging, monitoring, etc.

5.5 PoisonPill module

A simple class that is sent in a message to signal that the node should be terminated.

class `nanostream.message.poison_pill.`**PoisonPill**
> Bases: `object`

5.6 Trigger module

A simple class containing no data, which is intended merely as a trigger, signaling that the downstream node should do something.

class `nanostream.message.trigger.`**Trigger**(*previous_trigger_time=None*, *trigger_name=None*)
> Bases: `object`

`nanostream.message.trigger.`**hello_world**()

5.7 Batch module

We'll use markers to delimit batches of things, such as serialized files and that kind of thing.

class `nanostream.message.batch.`**BatchEnd**(**args*, ***kwargs*)
> Bases: `object`

class `nanostream.message.batch.`**BatchStart**(**args*, ***kwargs*)
> Bases: `object`

class `nanostream.message.canary.`**Canary**
> Bases: `object`

5.8 NanoStreamQueue module

These are queues that form the directed edges between nodes.

class `nanostream.node_queue.queue.`**NanoStreamQueue**(*max_queue_size*, *name=None*)
> Bases: `object`

---

empty

get()

put(*message*, **args*, *previous_message=None*, ***kwargs*)
    Places a message on the output queues. If the message is `None`, then the queue is skipped.

Messages are `NanoStreamMessage` objects; the payload of the message is message.message_content.

LICENSE

Copyright (C) 2016 Zachary Ernst zac.ernst@gmail.com

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <http://www.gnu.org/licenses/>.

INDICES AND TABLES

- genindex
- modindex
- search

PYTHON MODULE INDEX

n

H

I

K

L

M